

NEMO SYNTHESIS ENGINE - MARKETING GUIDE v2.0

Data Invisibility. Pure Intelligence. Your Complete Ownership.

EXECUTIVE SUMMARY

Nemo is the world's first truly private personal AI assistant.

Built on The Blanket Theory, Nemo synthesizes your keyboard behavior and screen context to predict what you'll do *before you do it*—without storing, recording, or transmitting a single byte of your data. Your data is not just private. It's invisible.

One program. 4 keys. Complete data ownership. Revolutionary prediction.

THE CORE PROMISE: DATA INVISIBILITY

What "Data Invisibility" Means

Your digital work disappears the moment it's analyzed.

- ✗ NOT: "We encrypt your data" (data still exists, just encrypted)
- ✗ NOT: "Your data is deleted after 30 days" (data exists temporarily)
- ✓ YES: "Your data never exists persistently" (analyzed in real-time, forgotten immediately)

How We Achieve It

Everything happens in RAM. Nothing ever touches disk.

- Keystroke → Analyzed → Synthesized → Forgotten
 - Voice Input → Transcribed → Converted to Text → Audio Buffer Cleared
 - Screen State → Analyzed → Context Extracted → Forgotten
 - AI Response → Spoken → Audio Memory Cleared
 - **Result: Zero Data Residue**
-

THE REVOLUTIONARY PROBLEM WE SOLVE

Why Current AI Sucks

They're all surveillance tools pretending to be assistants.

Current AI

NEMO

Records your keyboard	Analyzes patterns → forgets
Records your voice	Transcribes → clears immediately
Records your screen	Understands context → forgets
Stores everything forever	Stores nothing ever
Sells your data to advertisers	Can't—we don't have it
Requires constant internet	Works offline with Ollama
Treats you like everyone else	Learns YOUR unique style
You're the product	You own everything

NEMO reverses the entire paradigm.

THE FOUR-BUTTON REVOLUTION

Button 1: RIGHT SHIFT → Speech-to-Text + Reading

Your voice becomes your interface.

- **Press & Hold:** Start speaking
 - **Real-Time Transcription:** See your words appear as you speak
 - **Auto-Paste:** Release to insert transcribed text anywhere
 - **Read Highlighted Text:** Hold RIGHT SHIFT with text selected to hear it read aloud
 - **Use Cases:**
 - Dictate emails without typing
 - Write documents hands-free
 - Have text read to you while you work
 - No audio recording—just transcription
-

Button 2: RIGHT ALT → Gemini Voice AI + Screenshot Capture

Instant access to your personal AI butler with full context.

- **Press & Hold:** Start recording audio
 - **Live Transcription:** See what you're saying in real-time
 - **Screenshot Option (configurable):** RIGHT ALT automatically captures your current screen when pressed
 - **Send to Gemini:** Your voice + optional screenshot goes to Gemini API
 - **Voice Response:** Gemini's answer comes back as spoken audio
 - **Configure:** nemo config gemini to toggle screenshot/video capture on/off
 - **Use Cases:**
 - "What do I need to fix in this code?" (with screenshot)
 - "Summarize this spreadsheet" (with live data)
 - "What's the next step?" (Gemini sees your current context)
 - Get intelligent answers with full visual context
-

Button 3: RIGHT ALT + LEFT ARROW → Temporal Rewind

Navigate backwards through your work by inference.

- **Function:** Infer what you were doing 1-5 minutes ago
 - **How:** Analyzes keyboard signature + screen context
 - **Returns:** Natural language synthesis of past actions
 - **No Recording:** Pure behavioral synthesis—no video, no replay
 - **Use Cases:**
 - "What was I working on 3 minutes ago?"
 - Recover lost context
 - Trace your work history without recordings
-

Button 4: RIGHT ALT + RIGHT ARROW → FORWARD (Temporal Prediction) ☆ REVOLUTIONARY

The game-changer: Predict what you'll do next.

- **Function:** Predicts your next action based on patterns
 - **Power:** Uses 35-dimensional keystroke signature + AI synthesis
 - **Speed:** Sub-100ms prediction—faster than you can think
 - **Accuracy:** Learns from your behavior, gets smarter daily
 - **Output:** Next keystroke prediction, next command suggestion, workflow prediction
 - **Why It's Revolutionary:**
 - Current AI is *reactive* (responds to what you did)
 - NEMO FORWARD is *proactive* (predicts what you'll do)
 - This is the future of personal computing
 - **Use Cases:**
 - Autocomplete commands before you type them
 - Suggest the next file you'll need
 - Predict the next line of code
 - "What am I about to do?" answered before you do it
 - **This is what separates NEMO from everything else**
-

THE TECHNOLOGY: DATA INVISIBILITY ARCHITECTURE

The Synthesis Engine (69.2K LOC)

Real-time understanding without storage.

Screen Analyzer

- Analyzes current screen 60 times per second
- Extracts: Application, window content, visual state
- **Never persists** to disk
- Fresh analysis each time
- **Data Invisibility:** Screen state → analyzed → forgotten

Keyboard Synthesizer

- Builds 35-dimensional keystroke signature:
 - **12 Timing Dimensions:** Keystroke speed, intervals, rhythm
 - **8 Pressure Dimensions:** Key force, acceleration, decay
 - **10 Pattern Dimensions:** Bigrams, trigrams, sequences
 - **5 Intent Vectors:** Detected user intent from typing behavior
- Updates in real-time as you type
- **Never stored permanently**
- **Data Invisibility:** Keystroke → synthesized → forgotten

Temporal Inference Engine

- **REWIND:** Combines historical keyboard signature with current screen context to infer what happened before
- **FORWARD:** Uses keystroke patterns + contextual AI to predict next action
- No video recording required
- Pure behavioral synthesis
- **Data Invisibility:** Past/future inferred from patterns, patterns discarded

Multi-AI Integration

- **Gemini:** Google's advanced multimodal model (supports screenshots)
- **Claude:** Anthropic's reasoning engine
- **Ollama:** Open-source, runs locally (100% offline)
- **Custom Models:** Bring your own AI
- **Screenshot Feature:** Optional context capture for Gemini (configurable per mode)

Real-Time Processing

- All synthesis happens in-memory
- Sub-100ms response time
- Works on standard hardware
- Scales to any keystroke speed
- **No temporary files created**

Voice System (11.3K LOC)

Speech as your primary interface.

Speech-to-Text

- Captures voice input while RIGHT SHIFT held
- Real-time transcription as you speak
- Audio processed immediately
- Text extracted
- **Audio buffer cleared instantly**
- No temp files
- No audio logs

Text-to-Speech

- Converts any text to natural speech
- Multiple voices available (male/female/neutral)
- Adjustable speech rate
- **Audio plays in memory only**
- No written to disk
- No logging

Gemini Integration

- Right ALT hotkey sends: Voice input + optional screenshot
- Gemini processes and responds
- Response comes back as speech
- **Data Invisibility:** Request → processed externally → response returned → forgotten locally

Zero-Storage Architecture (13.0K LOC)

Guaranteed data invisibility.

In-Memory Processing Pipeline

Input → Buffer (RAM) → Process → Output → Buffer Cleared

Each stage:

1. Data enters RAM buffer
2. Analyzed/processed
3. Result extracted
4. Original data cleared from memory
5. **Zero persistence possible**

8-Point Security Audit

Every claim verified. Run anytime: `nemo security verify`

1. ✓ **Temp Directory Check:** No audio files, no keystroke logs
2. ✓ **Cache Directory Verification:** No persistence between runs
3. ✓ **Memory Forensics:** No strings left in process memory
4. ✓ **Log File Analysis:** No sensitive data in logs
5. ✓ **Credential Storage Check:** Encrypted, never logged
6. ✓ **Clipboard Monitoring:** Not captured or persisted
7. ✓ **Network Traffic Analysis:** No data exfiltration
8. ✓ **Behavioral Verification:** Synthesis only, no recording

WHY NEMO IS REVOLUTIONARY

1. Data Invisibility vs Everyone Else

Competitor	Storage	Record?	Data Stays?
------------	---------	---------	-------------

Microsoft Copilot	Yes	Always	Forever
-------------------	-----	--------	---------

ChatGPT	Yes	Logs	30 + days
Google Assistant	Yes	Audio files	Indefinite
Alexa	Yes	Continuous	"We delete it"
NEMO	Never	Never	Never

2. FORWARD Key is Unique

- **Rewind:** Infer past (interesting but only 50% useful)
- **FORWARD:** Predict future (game-changing and 90% useful)
- Every competitor has "history replay"
- **ONLY NEMO has true prediction**

3. Synthesis vs Recording

- Recording = Limited to what you did (playback only)
- Synthesis = Unlimited inference (understand intent, predict actions)
- Recording requires storage = privacy risk
- Synthesis requires no storage = data invisibility

4. Configurable Privacy

```
nemo config gemini
# Toggle: Screenshot capture (on/off)
# Toggle: Video recording (on/off per mode)
# Toggle: Cloud vs Local processing
# Control what Gemini can see
```

5. Offline Capable

- Install Ollama
- NEMO runs completely offline
- No API keys needed
- No internet required
- Full functionality, zero cloud dependency

6. Your Data, Your Rules

- No terms of service
- No data sharing
- No third-party access
- You own everything
- Forever free

THE TECHNICAL ADVANTAGE

Component	Status	LOC
Synthesis Engine	✓ Complete	69.2K
Voice System	✓ Complete	11.3K

Zero-Storage Architecture	✓ Complete 13.0K
Four-Button Interface	✓ Complete 8.5K
CLI System	✓ Complete 15.2K
Security Verification	✓ Complete 8.0K
Multi-AI Integration	✓ Complete 22.1K
RL Training Environments	✓ Complete 42.0K
TOTAL	✓ READY 239.3K

REAL-WORLD EXAMPLES

Example 1: Developer with RIGHT SHIFT

1. Hold RIGHT SHIFT
2. Say: "Create a login form in React"
3. Release RIGHT SHIFT
4. Text appears in editor automatically
5. No recording—just transcription

Example 2: Manager with RIGHT ALT + Screenshot

1. Stand in front of complex spreadsheet
2. Hold RIGHT ALT
3. Say: "Summarize this month's results"
4. Optional: Screenshot captures spreadsheet
5. Gemini analyzes and responds with spoken summary

Example 3: Analyst with FORWARD Key

1. Working on data analysis
 2. Press RIGHT ALT + RIGHT ARROW
 3. NEMO predicts: "You'll open the database query next"
 4. Autocomplete suggests correct SQL
 5. Save 10 seconds with FORWARD prediction
-

CONFIGURATION OPTIONS

Gemini Mode Settings

```
nemo config gemini --help
```

Options:

--screenshot ON/OFF	Enable/disable screenshot capture
--video ON/OFF	Enable/disable video recording
--cloud-only	Use Gemini API only
--hybrid	Combine local + cloud processing
--context-level	What to capture (minimal/normal/full)

RIGHT SHIFT Settings

```
nemo config speech-to-text --help
```

Options:

--auto-paste ON/OFF	Auto-insert transcribed text
--real-time-display ON/OFF	Show live transcription
--read-highlighted ON/OFF	Read selected text aloud
--timeout-seconds	How long to listen (default: 10)

FORWARD Key Settings

```
nemo config forward --help
```

Options:

--predictions 1-5	How many predictions to show
--confidence-threshold	Minimum confidence to suggest
--learning-mode ON/OFF	Learn from your corrections
--prediction-horizon	How far ahead to predict (1-30s)

SYSTEM REQUIREMENTS

Minimum (Offline Mode)

- Python 3.10 +
- 200MB disk space
- 100MB RAM
- Windows, macOS, or Linux
- Ollama (optional, for offline)

Recommended (Cloud + Gemini)

- Python 3.11 +
- 1GB disk space
- 300MB RAM
- SSD preferred
- 5MB/sec internet (for Gemini API)

INSTALLATION

1. Download from downloadnemo.com

```
# Visit https://downloadnemo.com
# Download Nemo v1.0.0
# Extract ZIP
# Run: pip install -r requirements.txt
```

2. Configure for Your Needs

```
nemo setup
# Choose AI model (Gemini/Claude/Ollama)
```

```
# Configure Gemini screenshot settings  
# Map your buttons  
# Test audio
```

3. Run and Enjoy

```
nemo buttons start  
# System ready  
# Press RIGHT SHIFT to speak  
# Press RIGHT ALT for Gemini  
# Press RIGHT ALT + RIGHT ARROW to predict next action
```

THE PROMISE: DATA INVISIBILITY GUARANTEE

We make one simple promise:

Your data never persists. Your work never lingers. Your thoughts remain yours.

This isn't marketing. It's verifiable. Run `nemo security verify` anytime and audit our claims.

PRICING

Plan	Cost	Includes
NEMO Core	Free	Everything listed above
NEMO + Gemini	Pay-as-you-go	Google Gemini API (Optional)
NEMO + Claude	Pay-as-you-go	Anthropic Claude API (Optional)
NEMO + Ollama	Free	Open-source models (Local)

Most users need \$0. Some use \$5-20/month for API costs.

COMPETITIVE ADVANTAGES

1. **Data Invisibility** - Only we achieve true zero-persistence
 2. **FORWARD Key** - Only predictive AI assistant (not just reactive)
 3. **Screenshot Integration** - Optional, configurable context for Gemini
 4. **Offline Capable** - Works without internet (with Ollama)
 5. **Synthesis vs Recording** - Intelligence, not surveillance
 6. **Open Architecture** - Audit our code anytime
 7. **Security Verified** - 8-point audit system
 8. **Completely Free** - No subscriptions, forever
 9. **Unique Hotkeys** - 4 buttons, infinite functionality
 10. **The Blanket Theory** - Philosophy that respects human dignity
-

ROADMAP

v1.0.0 (Current) ✓

- ✓ Synthesis engine
- ✓ Four-button interface
- ✓ Speech-to-text with RIGHT SHIFT
- ✓ Gemini with screenshot capture
- ✓ FORWARD key (temporal prediction)
- ✓ Zero-storage guarantee
- ✓ Data invisibility verified

v2.0 (Next)

- Advanced FORWARD predictions (10+ seconds ahead)
- Video recording option for Gemini (optional, configurable)
- Multi-window synthesis
- Team collaboration mode (privacy-preserving)
- Custom keyboard layouts

v3.0 (Vision)

- NEMO becomes your second brain
 - Predicts problems before they happen
 - Understands your entire digital life
 - Perfect personal AI assistant
 - Your work history reconstructed from synthesis alone (no recordings)
-

FAQ

Q: My data is really never stored? A: Never. Run `nemo security verify`. 8-point audit confirms it. Open source our code and verify yourself.

Q: What about Gemini API? Does Google keep my data? A: You control that via `nemo config gemini`. Toggle screenshot OFF if you don't want Gemini seeing your screen. Claude and Ollama are alternatives.

Q: Is FORWARD really predictive? A: Yes. It predicts your next keystroke/action using your unique 35-dimensional behavioral signature. Gets smarter as you use it.

Q: Can I configure which buttons do what? A: Completely. `nemo config buttons` lets you remap everything.

Q: Does this work offline? A: Core NEMO works fully offline. FORWARD key works offline. Gemini requires internet (unless using Ollama locally).

Q: Why is this free? A: We believe AI should serve humans, not exploit them. You own your data. That's the whole point.

Q: What's the catch? A: No catch. We're funded by people who believe in data invisibility. That's it.

GET STARTED TODAY

1. Visit: <https://downloadnemo.com>
2. Download: Nemo v1.0.0
3. Run: nemo setup
4. Press: RIGHT SHIFT to speak
5. Experience: The future of personal AI

NEMO: Data Invisibility. Pure Intelligence. Your Future.

Made in America. For humans. By people who respect privacy.

SUPPORT

- **GitHub:** <https://github.com/torresjchristopher/nemo>
 - **Issues:** Report bugs and feature requests
 - **Security:** Responsible disclosure welcome
 - **Community:** Share your NEMO workflows
-

NEMO: The revolutionary personal AI that respects data invisibility.

Your data. Your rules. Your future.